

Transactional recovery read — stall cheat sheet

KafkaPartitionPersistence · read_committed drain to the captured high-watermark · poll 10 ms · deadline 3 min · legit wait ~70 s (txn.timeout 60 s + abort scan 10 s) · worked example HW = 1000

CONTROL FLOW one loop, two kinds of turn, two exits

SETUP (once, before the loop)
BP1 highWatermark = endOffsets(read_uncommitted) · the high watermark
BP2 lastStableOffset = endOffsets(read_committed) · the LSO (reader's own end)
BP3 if lastStableOffset < highWatermark → WARN "recovery waits for open transaction(s)"
BP4 targetOffset = highWatermark · the target: captured once, fixed for the read (never the LSO)

LOOP (tailRecM · one turn ≈ 10 ms)
BP5 offset = position()
└ if offset ≥ targetOffset → EXIT ✓ read complete
BP6 advanced? (offset ≠ last.position)
└ YES → ADVANCING turn · reset clock (progressAt = now)
└ NO → WAITING turn · BP7 if now – progressAt ≥ 3 min → EXIT ✗ BP8 RecoveryReadStalledError
BP9 logEnd < target ? truncation : outlived
else · every 5 s → INFO "no progress ..."
↳ both turns end: poll(10.millis) ; fold batch → back to BP5

TIMELINE — FOUR SCENARIOS SIDE BY SIDE a walk through the loop above

target = high watermark, captured once at BP1/BP4 and held constant · ♻ = the loop, one per ~10 ms turn · advancing turn (BP6) poll → batch, position climbs · waiting turn (BP7) poll → empty, position frozen, clock ticks

wall clock	A — no wait (takeover-abort)	B — awaiting open txn (resolves)	C — hanging transaction	D — truncation
0 entry	BP1–BP4: LSO 1000 = HW → no warn, target 1000 (own txn aborted at takeover)	BP1–BP4: LSO 950 < HW 1000 → WARN, target 1000	same: LSO 950 → WARN, target 1000	LSO 1000 = HW → no warn, target 1000
0–~50 ms climb	♻ advancing turns (BP6) · position 0→1000 → BP5 exit ✓ complete @1000 (only advancing turns — never BP7)	♻ advancing turns (BP6) · position 0→950, each poll a batch	♻ advancing turns (BP6) · position 0→950	♻ advancing turns (BP6) · climbing... ⚡ election → log end 900
first stall (no log yet)	— done —	waiting turn (BP7) begins · stuck@950, stalledFor=0	waiting turn (BP7) · stuck@950	waiting turn (BP7) · stuck@900 (nothing past 900)
5 s → 70 s steady loop	— done —	♻ waiting turns (BP7) ≈500×/5 s · poll(10ms) empty, position frozen, BP8 false · every 5 s → INFO "... Ns"	same @950	same @900
≈70 s	— done —	⚡ txn resolves → advancing turn: poll returns 950–999 → position 1000 → BP5 exit ✓ complete @1000	still waiting @950 (♻ INFO every 5 s, BP8 false)	still waiting @900 (♻ INFO every 5 s)
180 s deadline	— done —	— done —	waiting turn · stalledFor≥3min → BP8 ✗ → BP9 logEnd 1000 ≥ target → "outlived deadline" → RecoveryReadStalledError@950	→ BP8 ✗ → BP9 logEnd 900 < target → "truncation" → ...@900

WHAT YOU SEE IN THE LOGS verbatim; blue = interpolated values

at	level	message
BP3 start	WARN	Snapshot topic <t> partition <p> recovery waits for open transaction(s): last-stable-offset <LS0> below target <HW>; normally resolves within the pinning producer's transaction.timeout.ms plus the broker's abort scan
BP7 waiting	INFO	Snapshot topic read making no progress at offset <o>, target <t>, stalled for <N>s (repeats every 5 s)
BP8 fail	ERROR	recovery read of <tp> made no progress at offset <pos>, short of target <target>, past the stall deadline - failing rather than hanging or silently recovering possibly incomplete state; <diagnosis>
BP9 ↳ trunc.	diag	the log end regressed to <logEnd>, below the captured target: log truncation after an unclean leader election - acknowledged snapshot records are lost
BP9 ↳ outlived	diag	the log end <logEnd> still covers the target: an open transaction outlived the deadline - ... detect and abort it with kafka-transactions.sh

Every scenario runs the same loop — each turn is either an advancing turn (BP6: poll returned a batch, position climbs) or a waiting turn (BP7: poll came back empty, position frozen, clock ticks). A is all advancing turns; B/C/D add a run of waiting turns. B ≡ C until 70 s — identical waiting-turn runs; the only difference is whether a poll returns records before the clock reaches BP8. No separate "hang" detector — the deadline is it, discovered late on an ordinary turn, never a scheduled wake-up. C vs D differ only at BP3 (warn fires / silent) and BP9 (LogEnd vs target). Reader is group-less, so max.poll.interval.ms eviction can't unwind the loop — only ✓ BP5 (records) or ✗ BP8 (deadline) exits it.

CAVEATS & OPERATIONS truncation is a Kafka durability boundary, not a bug

Two stalls, two shapes — B/C stall at the LSO: committed records sit above, hidden by an open transaction (they appear once it resolves). D stalls at the true log end: those records are gone (truncated). Detection is in-flight only — the deadline catches truncation that lands during a read. A truncation between reads is silent: the next recovery adopts the regressed end as its target and completes, possibly resuming stale state from an ahead offset — the #732 shape. Reduce the loss risk — the deadline only flags lost records; it can't recover them. Keep the snapshot topic durable: unclean.leader.election.enable=false, min.insync.replicas ≥ 2, RF ≥ 3 (the transactional producer already forces acks=all). On RecoveryReadStalledError — truncation → offset-reset / restore (don't blind-restart: it recovers silently to the regressed end); outlived → wait, or abort a hanging txn with kafka-transactions.sh.